

Language modeling Seq2seq

NLP

lecturer: Mollaev D. E.
Sber AI Lab

Lecture plan

- What is the language modeling?
- General framework
- Neural LMs
- Generating strategy
- Evaluation LMs
- Seq2seq framework
- Attention

Lecture plan

- What is the language modeling?
- General framework
- Neural LMs
- Generating strategy
- Seq2seq framework
- Attention

Train Models

- have some properties of trains (look like ones)
- can behave similarly to trains
- good models have more of the above




Models of Physical World


- understand which events are in better agreement with the world, which are more likely
- can predict what happens given some “context”




Will it fall?



In which direction?



Different masses



Complex scenes




Infer the mass



Predict fluids

Language Models



Language Models

The intuition is exactly the same!

What is different, is the notion of an event: for language, an event is a linguistic unit (text, sentence, token, symbol).

Language Models (LMs) estimate the probability of different linguistic units: symbols, tokens, token sequences.

We do we need it?

We deal with Language
Models every day!

Web search engine / ...

I saw a cat|

I saw a cat on the chair

I saw a cat running after a dog

I saw a cat in my dream

I saw a cat book

We do we need it?

We deal with Language Models every day!

Translation service / mail agent / ...

I saw a ca|

car ←

We do we need it?

We deal with Language
Models every day!

Translation service / mail agent / ...

I saw a catt

Probably you meant I saw a cat

We do we need it?

We deal with Language
Models every day!

Keyboard / mail agent / ...

I saw a catt

cat
car


Ambiguity



Similarly sounding options

She studies morphosyntax
She studies more faux syntax
She studies morph or syntax

....




The morphosyntax example is from the slides by Alex Lascarides and Sharon Goldwater, Foundations of Natural Language Processing course at the University of Edinburgh.

Lecture plan

- What is the language modeling?
- General framework
- Neural LMs
- Generating strategy
- Seq2seq framework
- Attention


Modeling

What is the probability
to pick a green ball?



Modeling


What is the probability
to pick a green ball?



$$\frac{5}{5 + 6 + 4 + 3} = \frac{5}{18}$$


Modeling

What is the probability to pick a green ball?



$$\frac{5}{5 + 6 + 4 + 3} = \frac{5}{18}$$

Can we do the same for sentences?




$$P(\text{the mat is timming the tebn}) = \frac{0}{|\text{corpus}|} = 0$$

$$P(\text{mat the timming tebn is the}) = \frac{0}{|\text{corpus}|} = 0$$

With this approach, sentences that never occurred in the corpus will receive zero probability


Modeling

What is the probability to pick a green ball?



$$\frac{5}{5 + 6 + 4 + 3} = \frac{5}{18}$$

Can we do the same for sentences?



$$P(\text{the mat is timming the tebn}) = \frac{0}{|\text{corpus}|} = 0$$

$$P(\text{mat the timming tebn is the}) = \frac{0}{|\text{corpus}|} = 0$$


With this approach, sentences that never occurred in the corpus will receive zero probability

But the first sentence is “more likely” than the second!
This method is not good!




Modeling

What is the probability to pick a green ball?



$$\frac{5}{5 + 6 + 4 + 3} = \frac{5}{18}$$

Can we do the same for sentences?



$$P(\text{the mat is timming the tebn}) = \frac{0}{|\text{corpus}|} = 0$$

$$P(\text{mat the timming tebn is the}) = \frac{0}{|\text{corpus}|} = 0$$

With this approach, sentences that never occurred in the corpus will receive zero probability

But the first sentence is “more likely” than the second!
This method is not good!



We can not estimate sentence probabilities reliably if we treat them as atomic units!

Modeling

Image we

- read the sentence **I saw a cat on a mat** word by word,
- update probability every time we see a new token

$$P(\mathbf{I}) =$$

$$P(\mathbf{I})$$



Probability of \mathbf{I}

Modeling

Formally,

- (y_1, y_2, \dots, y_n) is a sequence of tokens,
- $P(y_1, y_2, \dots, y_n)$ - probability to see these tokens (in this order)

Using the product rule of probability (aka “chain rule”), we get:

$$P(y_1, y_2, \dots, y_n) = P(y_1) \cdot P(y_2|y_1) \cdot P(y_3|y_1, y_2) \cdot \dots \cdot P(y_n|y_1, \dots, y_{n-1}) = \prod_{t=1}^n P(y_t|y_{<t})$$

Lecture plan

- What is the language modeling?
- General framework
- Neural LMs
- Generating strategy
- Seq2seq framework
- Attention

Neural LMs

$$P(y_1, y_2, \dots, y_n) = P(y_1) \cdot P(y_2|y_1) \cdot P(y_3|y_1, y_2) \cdot \dots \cdot P(y_n|y_1, \dots, y_{n-1}) = \prod_{t=1}^n P(y_t|y_{<t}).$$

We need to: define how to compute the conditional probabilities $P(y_t|y_1, \dots, y_{t-1})$.

How: Train a neural network to predict them.

Neural LMs


Intuitively, neural Language Models do two things:

- **process context → model-specific**

The main idea here is to get a vector representation for the previous context. Using this representation, a model predicts a probability distribution for the next token. This part could be different depending on model architecture (e.g., RNN, CNN, whatever you want), but the main point is the same - to encode context.


- **generate a probability distribution for the next token → model-agnostic**

Once a context has been encoded, usually the probability distribution is generated in the same way - see below.




Neural LMs: High-level pipeline

- feed word embedding for previous (context) words into a network;
- get vector representation of context from the network;
- from this vector representation, predict a probability distribution for the next token.



Neural LMs: Training



Neural LMs: Training

Formally, let's assume we have a training instance with the source $x = (x_1, \dots, x_m)$ and the target $y = (y_1, \dots, y_n)$. Then at the timestep t , a model predicts a probability distribution $p^{(t)} = p(\cdot|y_1, \dots, y_{t-1}, x_1, \dots, x_m)$. The target at this step is $p^* = \text{one-hot}(y_t)$, i.e., we want a model to assign probability 1 to the correct token, y_t , and zero to the rest.

The standard loss function is the cross-entropy loss. Cross-entropy loss for the target distribution p^* and the predicted distribution p is

$$\text{Loss}(p^*, p) = -p^* \log(p) = -\sum_{i=1}^{|V|} p_i^* \log(p_i).$$

Since only one of p_i^* is non-zero (for the correct token y_t), we will get


$$\text{Loss}(p^*, p) = -\log(p_{y_t}) = -\log(p(y_t|y_{<t}, x)).$$

Lecture plan

- What is the language modeling?
- General framework
- Neural LMs
- Generating strategy
- Seq2seq framework
- Attention

Generating strategy

As we saw before, to generate a text using a language model you just need to sample tokens from the probability distribution predicted by a model.



Generating strategy

Coherence and Diversity

You can modify the distributions predicted by a model in different ways to generate texts with some properties. While the specific desired text properties may depend on the task you care about (as always), usually you would want the generated texts to be:

- **coherent** - the generated text has to make sense;
- **diverse** - the model has to be able to produce very different samples.


Generating strategy

- Standard Sampling
- Sampling with temperature
 - Top-k sampling
 - Nucleus Sampling

Sampling with temperature

$$\frac{\exp(h^T w)}{\sum_{w_i \in V} \exp(h^T w_i)} \rightarrow \frac{\exp\left(\frac{h^T w}{\tau}\right)}{\sum_{w_i \in V} \exp\left(\frac{h^T w_i}{\tau}\right)}$$

τ - softmax temperature




Top-k sampling


Fixed K is not always good

While usually top-K sampling is much more effective than changing the softmax temperature alone, the fixed value of K is surely not optimal. Look at the illustration below.

Top-K for a flat distribution: not enough



Top-K for a peaky distribution: too many





The fixed value of K in the top-K sampling is not good because top-K most probable tokens may

- cover very small part of the total probability mass (in flat distributions);
- contain very unlikely tokens (in peaky distributions).

Nucleus Sampling

A more reasonable strategy is to consider not top-K most probable tokens, but top-p% of the probability mass: this solution is called **Nucleus sampling**.

Look at the illustration: with nucleus sampling, the number of tokens we sample from is dynamic and depends on the properties of the distribution.



Lecture plan

- What is the language modeling?
- General framework
- Neural LMs
- Generating strategy
- Evaluation LMs
- Seq2seq framework
- Attention

Perplexity and Text Log-Likelihood

$$L(y_{1:M}) = L(y_1, y_2, \dots, y_M) = \sum_{t=1}^M \log_2 p(y_t | y_{<t}) \quad Loss(y_{1:M}) = -\sum_{t=1}^M \log p(y_t | y_{<t})$$

Log-likelihood of the text

Note: cross-entropy (our loss)
is negative log-likelihood

It is more common to report its transformation called perplexity:

$$Perplexity(y_{1:M}) = 2^{-\frac{1}{M}L(y_{1:M})}$$

Low perplexity \rightarrow high data likelihood \rightarrow good!



Perplexity: Which Values Does it Take?

$$\text{Perplexity}(y_{1:M}) = 2^{-\frac{1}{M}L(y_{1:M})}$$

- The **best** perplexity is 1

If the model is perfect and assigns probability 1 to correct tokens, then the log-probabilities are zero

- The **worst** perplexity is $|V|$

If the model knows nothing about the data, it assigns probability $1/|V|$ to all tokens, regardless of context. Then:


$$\text{Perplexity}(y_{1:M}) = 2^{-\frac{1}{M}L(y_{1:M})} = 2^{-\frac{1}{M}\sum_{t=1}^M \log_2 p(y_t | y_{1:t-1})} = 2^{-\frac{1}{M} \cdot M \cdot \log_2 \frac{1}{|V|}} = 2^{\log_2 |V|} = |V|$$

Lecture plan

- What is the language modeling?
- General framework
- Neural LMs
- Generating strategy
- Seq2seq framework
- Attention

Machine Translation

- Translation between natural languages
- More generally, translation between any sequences



Machine Translation

Human Translation

$$y^* = \arg \max_y p(y|x)$$



The “probability” is
intuitive and is given
by a human
translator’s expertise

Machine Translation

Human Translation

$$y^* = \arg \max_y p(y|x)$$



The “probability” is
intuitive and is given
by a human
translator’s expertise

Machine Translation

$$y' = \arg \max_y p(y|x, \theta)$$

model

parameters

Machine Translation

Human Translation

$$y^* = \arg \max_y p(y|x)$$

The “probability” is
intuitive and is given
by a human
translator’s expertise

Machine Translation

$$y' = \arg \max_y p(y|x, \theta)$$

Questions we need to answer

- **modeling**

How does the model
for $p(y|x, \theta)$ look like?

Machine Translation

Human Translation

$$y^* = \arg \max_y p(y|x)$$

The “probability” is intuitive and is given by a human translator’s expertise

Machine Translation

$$y' = \arg \max_y p(y|x, \theta)$$

Questions we need to answer

- **modeling**

How does the model for $p(y|x, \theta)$ look like?

- **learning**

How to find θ ?

Machine Translation

Human Translation

$$y^* = \arg \max_y p(y|x)$$

The “probability” is intuitive and is given by a human translator’s expertise

Machine Translation

$$y' = \arg \max_y p(y|x, \theta)$$

Questions we need to answer

- **modeling**

How does the model for $p(y|x, \theta)$ look like?

- **learning**

How to find θ ?


- **search**

How to find the argmax?

Encoder-Decoder Framework

The standard modeling paradigm:


- **Encoder** – reads the source sentence and produces its representation



Encoder-Decoder Framework

The standard modeling paradigm:

- **Encoder** – reads the source sentence and produces its representation
- **Decoder** - uses source representation from the encoder to generate the target sequence.



Conditional Language Models

Language Models: $P(y_1, y_2, \dots, y_n) = \prod_{t=1}^n p(y_t | y_{<t})$
(left-to-right)

Conditional Language Models

Language Models: $P(y_1, y_2, \dots, y_n) = \prod_{t=1}^n p(y_t | y_{<t})$
(left-to-right)

Conditional
Language Models: $P(y_1, y_2, \dots, y_n, | \textcolor{brown}{x}) = \prod_{t=1}^n p(y_t | y_{<t}, \textcolor{brown}{x})$

condition on source x

General View


- process context – model-specific

Get vector representation
of the source and previous
target tokens


- evaluate probabilities –
model-agnostic

Predict probability
distribution for the
next target token


$$P(* | \text{Я видел котю на мате } \langle \text{eos} \rangle)$$




High-Level Pipeline



High-Level Pipeline



Two RNN Model



Generating

$$y' = \arg \max_y p(y|x) = \arg \max_y \prod_{t=1} p(y_t|y_{<t}, x)$$

Generating. Greed

$$y' = \arg \max_y p(y|x) = \arg \max_y \prod_{t=1} p(y_t|y_{<t}, x)$$

Straightforward:

- greedy - at each step, pick token with the highest probability

Generating. Greed Bad?

$$y' = \arg \max_y p(y|x) = \arg \max_y \prod_{t=1}^n p(y_t|y_{<t}, x)$$

Straightforward:

- greedy - at each step, pick token with the highest probability

$$\arg \max_y \prod_{t=1}^n p(y_t|y_{<t}, x) \neq \prod_{t=1}^n \arg \max_{y_t} p(y_t|y_{<t}, x)$$

- this is bad!

Beam Search

<bos>

Start with the begin of sentence token or with an empty sequence


Lecture plan

- What is the language modeling?
- General framework
- Neural LMs
- Generating strategy
- Seq2seq framework
- Attention

Attention


The Problem of Fixed Encoder Representation

Problem: Fixed source representation is suboptimal: (i) for the encoder, it is hard to compress the sentence; (ii) for the decoder, different information may be relevant at different steps.




Attention: A High-Level View


Attention: At different steps, let a model "focus" on different parts of the input.



Attention



How to Compute Attention Score?



Dot-product

$$h_t^T \times \begin{bmatrix} \textcolor{green}{\circ} \\ \textcolor{green}{\circ} \\ \textcolor{green}{\circ} \\ \textcolor{green}{\circ} \end{bmatrix} s_k$$

$$\text{score}(h_t, s_k) = h_t^T s_k$$

Bilinear

$$h_t^T \times \begin{bmatrix} W \end{bmatrix} \times \begin{bmatrix} \textcolor{green}{\circ} \\ \textcolor{green}{\circ} \\ \textcolor{green}{\circ} \\ \textcolor{green}{\circ} \end{bmatrix} s_k$$

$$\text{score}(h_t, s_k) = h_t^T W s_k$$

Multi-Layer Perceptron

$$\textcolor{blue}{W}_2^T \times \tanh \left(\begin{bmatrix} W_1 \end{bmatrix} \times \begin{bmatrix} \textcolor{red}{\circ} \\ \textcolor{red}{\circ} \\ \textcolor{red}{\circ} \\ \textcolor{red}{\circ} \end{bmatrix} h_t \right) \times \begin{bmatrix} \textcolor{green}{\circ} \\ \textcolor{green}{\circ} \\ \textcolor{green}{\circ} \\ \textcolor{green}{\circ} \end{bmatrix} s_k$$


$$\text{score}(h_t, s_k) = w_2^T \cdot \tanh(W_1[h_t, s_k])$$

Attention

L'accord sur la zone économique européenne a été signé en août 1992.

The agreement on the European Economic Area was signed in August 1992.

<end>



Il convient de noter que l'environnement marin est le moins connu de l'environnement.

It should be noted that the marine environment is the least known of environments.

<end>

