

Transformer

NLP

lecturer: Mollaev D. E.
Sber AI Lab


Lecture plan

- Tokenization
- General view:
 - Seq2Seq
 - Ablation: transformer vs rnn with attention
- Transformer: architectures
 - Self-Attention module
 - Feed forward module
 - Positional encoding
 - Other important parts
- Training Details
- BERT-like
- GPT-like

Lecture plan


- Tokenization
- General view:
 - Seq2Seq
 - Ablation: transformer vs rnn with attention
- Transformer: architectures
 - Self-Attention module
 - Feed forward module
 - Positional encoding
 - Other important parts
- Training Details
- BERT-like
- GPT-like

Tokenization



Tokenization: BPE - training

From individual symbols, whole words and N-grams – to most common co-occurrences



To build a tokenizer vocabulary:

- Initialize a vocabulary with all unique symbols from a dataset
- Build a character co-occurrences model
- Pick the most common pair
- Repeat...

Tokenization: BPE - inference

The algorithm starts with segmenting a word into a sequence of characters. After that, it iteratively makes the following two steps until no merge is possible:

- among all possible merges at this step, find the highest merge in the table;
- apply this merge.

u-n-r-e-l-a-t-e-d
u-n re-l-a-t-e-d
u-n re-l-at-e-d
u-n re-l-ate-d
un re-l-at-ed
un re-l-ated
un rel-ated
un related
un@@ related

No merges
possible -> end

Lecture plan

- Tokenization
- General view:
 - Seq2Seq
 - Ablation: transformer vs rnn with attention
- Transformer: architectures
 - Self-Attention module
 - Feed forward module
 - Positional encoding
 - Other important parts
- Training Details
- BERT-like
- GPT-like

General View


- process context – model-specific

Get vector representation
of the source and previous
target tokens

- evaluate probabilities –
model-agnostic

Predict probability
distribution for the
next target token

$$P(* | \text{Я видел котю на мате } \langle \text{eos} \rangle)$$




Transformer: Attention is All You Need

	Seq2seq without attention	Seq2seq with attention	Transformer
processing within encoder	RNN/CNN	RNN/CNN	attention
processing within decoder	RNN/CNN	RNN/CNN	attention
decoder-encoder interaction	static fixed- sized vector	attention	attention


Lecture plan

- Tokenization
- General view:
 - Seq2Seq
 - Ablation: transformer vs rnn with attention
- Transformer: architectures
 - Self-Attention module
 - Feed forward module
 - Positional encoding
 - Other important parts
- Training Details
- BERT-like
- GPT-like


Transformer: overview



Transformer: encoder, decoder



Transformer: encoder blocks



Lecture plan

- Tokenization
- General view:
 - Seq2Seq
 - Ablation: transformer vs rnn with attention
- Transformer: architectures
 - Self-Attention module 
 - Feed forward module
 - Positional encoding
 - Other important parts
- Training Details
- BERT-like
- GPT-like

Transformer: self-attention

Each vector receives three representations (“roles”)

$$[W_Q] \times [] = []$$

Query: vector from which the attention is looking

“Hey there, do you have this information?”

$$[W_K] \times [] = []$$


Key: vector at which the query looks to compute weights

“Hi, I have this information – give me a large weight!”

$$[W_V] \times [] = []$$

Value: their weighted sum is attention output

“Here’s the information I have!”




- **query** - asking for information
- **key** - saying that it has some information
- **value** - giving the information

Transformer: self-attention


$$\text{Attention}(q, k, v) = \underbrace{\text{softmax}\left(\frac{qk^T}{\sqrt{d_k}}\right)v}_{\substack{\text{Attention weights} \\ \text{vector dimensionality of } K, V}}$$

from to

Transformer: self-attention




Transformer: self-attention



Multiplying x_1 by the W^Q weight matrix produces q_1 , the "query" vector associated with that word. We end up creating a "query", a "key", and a "value" projection of each word in the input sentence.


Transformer: self-attention





- For a vector \mathbf{X} , compute \mathbf{Q} , \mathbf{K} , and \mathbf{V} by multiplying learnable matrices $\mathbf{W}_{Q/K/V}$
- For a fixed q_i , compute attention **score** by matching it against \mathbf{K}
- Assemble result: retrieve a **value** with respect to its **score**

Transformer: masked self-attention

To forbid the decoder to look ahead, the model uses masked self-attention: future tokens are masked out.



Transformer: multi-head self-attention




Lecture plan

- Tokenization
- General view:
 - Seq2Seq
 - Ablation: transformer vs rnn with attention
- Transformer: architectures
 - Self-Attention module
 - Feed forward module ←
 - Positional encoding
 - Other important parts
- Training Details
- BERT-like
- GPT-like

Transformer: Feedforward block

$$FFN(x) = \max(0, xW_1 + b_1)W_2 + b_2.$$


- **Attention** - "look at other tokens and gather information"
- **FFN** - "take a moment to think and process this information"



Lecture plan

- Tokenization
- General view:
 - Seq2Seq
 - Ablation: transformer vs rnn with attention
- Transformer: architectures
 - Self-Attention module
 - Feed forward module
 - Positional encoding ←
 - Other important parts
- Training Details
- BERT-like
- GPT-like

Transformer: positional encoding




Lecture plan

- Tokenization
- General view:
 - Seq2Seq
 - Ablation: transformer vs rnn with attention
- Transformer: architectures
 - Self-Attention module
 - Feed forward module
 - Positional encoding
 - Other important parts ←—
- Training Details
- BERT-like
- GPT-like


Transformer: Other important parts

Residual connection




Residual connection:
add a block's input to
its output

Layer norm




Transformer: layer norm



- **Layer norm:** По каждомуциальному примеру: нормализует значения внутри одного слоя.
- **Batch norm:** По мини-батчу: для каждого признака вычисляет среднее и стандартное отклонение по всему батчу.

Transformer: layer norm




Оригинальная версия из *Attention Is All You Need* использует **Post-LN**, однако современные LLM без исключения строятся на **Pre-LN**. Разница в их поведении заключается в следующем:


- **Post-LN** показывает хорошие результаты на небольших трансформерах (до 6 слоёв), но при увеличении глубины обучения становится нестабильным и может расходиться.
- **Pre-LN** обеспечивает более стабильное обучение, что делает его предпочтительным для глубоких моделей.

Decoder

- Cross-Attention: **K** and **V** are coming from encoder, **Q** is from previous layer
- Masked MHA: assign large negative number to any token from the future



Transformer: overview



Lecture plan


- Tokenization
- General view:
 - Seq2Seq
 - Ablation: transformer vs rnn with attention
- Transformer: architectures
 - Self-Attention module
 - Feed forward module
 - Positional encoding
 - Other important parts
- Training Details and Hyperparameters:
- BERT-like
- GPT-like

Training details and hyperparameters

- Cross-Entropy Loss $NLL(y_{1:M}) = - \sum_{t=1}^M \log p(y_t|t_{t-1})$
- Teacher Forcing for decoder
- Adam Optimizer
- lr with warm-up scheduling $lr = d_{model}^{-0.5} \cdot \min(step_num^{-0.5}, step_num \cdot warmup^{-1.5})$
- BPE

Hypers:

- n_blocks (6 blocks)
- n_head (8 heads)
- n_dim (512)




Lecture plan


- Tokenization
- General view:
 - Seq2Seq
 - Ablation: transformer vs rnn with attention
- Transformer: architectures
 - Self-Attention module
 - Feed forward module
 - Positional encoding
 - Other important parts
- Training Details
- BERT-like
- GPT-like

BERT vs GPT

- Transformer (invented for translation task):



BERT (Bidirectional Encoder Representations from Transformers): overview



Model architecture:

- Transformer's encoder

What is special about it:

- Training objectives
 - MLM: Masked language modeling
 - NSP: Next sentence prediction
- The way it is used
 - No task-specific models


BERT - Training input

Training Input: Pairs of Sentences with Special Tokens

[CLS]: Special token


- Training time: predict if sentences are consecutive or not (Next Sentence Prediction /NSP objective)
 - Test time: downstream tasks (e.g., classification)

[SEP]: Special token-separator



Training on pairs of sentences: either consecutive or random (50%/50%)

BERT: cls token



BERT: pre-training objectives NSP

Input: [CLS] the man went to [MASK] store [SEP] he bought a gallon [MASK] milk [SEP]

Label: isNext


Input: [CLS] the man went to [MASK] store [SEP] penguin [MASK] are flight ##less birds [SEP]

Label: notNext

BERT: pre-training objectives MLM

- select some tokens
(each token is selected with the probability of 15%)
- replace these selected tokens
(with the special token [MASK] - with p=80%, with a random token - with p=10%, with the original token (remain unchanged) - with p=10%)
- predict original tokens (compute loss).

BERT: pre-training objectives MLM







At each training step:

- pick randomly 15% of tokens
- replace each of the chosen tokens with something
- predict original chosen tokens

BERT: fine-tuning


- Single sentence classification
- Sentence pair classification
- Question Answering
- Single sentence tagging (NER)



Lecture plan

- Tokenization
- General view:
 - Seq2Seq
 - Ablation: transformer vs rnn with attention
- Transformer: architectures
 - Self-Attention module
 - Feed forward module
 - Positional encoding
 - Other important parts
- Training Details
- BERT-like
- GPT-like

GPT-like



- GPT-1: Improving Language Understanding by Generative Pre-Training
- GPT-2: Language Models are Unsupervised Multitask Learners
- GPT-3: Language Models are Few-Shot Learners